

- Carla E Brodley and Mark A Friedl. Identifying mislabeled training data. *Journal of artificial intelligence research*, 11:131–167, 1999.
- Haw-Shiuan Chang, Erik Learned-Miller, and Andrew McCallum. Active Bias: Training More Accurate Neural Networks by Emphasizing High Variance Samples. In *Advances in Neural Information Processing Systems*, pp. 1002–1012, 2017.
- Pratik Chaudhari, Anna Choromanska, Stefano Soatto, Yann LeCun, Carlo Baldassi, Christian Borgs, Jennifer Chayes, Levent Sagun, and Riccardo Zecchina. Entropy-SGD: Biasing Gradient Descent Into Wide Valleys. *ICLR '17*, 2016.
- Terrance DeVries and Graham W Taylor. Improved regularization of convolutional neural networks with cutout. *arXiv preprint arXiv:1708.04552*, 2017.
- Yang Fan, Fei Tian, Tao Qin, and Jiang Bian. Learning What Data to Learn. *arXiv preprint arXiv:1702.08635*, 2017.
- Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *Proc. of ICML*, 2017.
- S. Hochreiter and J. Schmidhuber. Flat minima. *Neural Computation*, 9(1):1–42, 1997.
- Gao Huang, Zhuang Liu, Laurens van der Maaten, and Kilian Q Weinberger. Densely Connected Convolutional Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 4700–4708, 2017.
- Lu Jiang, Zhengyuan Zhou, Thomas Leung, Li-Jia Li, and Li Fei-Fei. MentorNet: Learning data-driven curriculum for very deep neural networks on corrupted labels. In *Proceedings of the 35th International Conference on Machine Learning*. PMLR, 2018.
- George H John. Robust decision trees: removing outliers from databases. In *Proceedings of the First International Conference on Knowledge Discovery and Data Mining*, pp. 174–179. AAAI Press, 1995.
- Angelos Katharopoulos and François Fleuret. Not all samples are created equal: Deep learning with importance sampling. In Jennifer G. Dy and Andreas Krause (eds.), *ICML*, volume 80 of *JMLR Workshop and Conference Proceedings*, pp. 2530–2539. JMLR.org, 2018. URL <http://dblp.uni-trier.de/db/conf/icml/icml2018.html#KatharopoulosF18>.
- Nitish Shirish Keskar, Dheevatsa Mudigere, Jorge Nocedal, Mikhail Smelyanskiy, and Ping Tak Peter Tang. On large-batch training for deep learning: Generalization gap and sharp minima. *arXiv preprint arXiv:1609.04836*, 2016.
- Tae-Hoon Kim and Jonghyun Choi. Screenernet: Learning curriculum for neural networks. *CoRR*, abs/1801.00904, 2018. URL <http://dblp.uni-trier.de/db/journals/corr/corr1801.html#abs-1801-00904>.
- Diederik Kingma and Jimmy Ba. Adam: A method for stochastic optimization, 2014. URL <http://arxiv.org/abs/1412.6980>. cite arxiv:1412.6980Comment: Published as a conference paper at the 3rd International Conference for Learning Representations, San Diego, 2015.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, and Others. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, pp. 201611835, 2017.
- Robert Kleinberg, Yuanzhi Li, and Yang Yuan. An alternative view: When does sgd escape local minima? *CoRR*, abs/1802.06175, 2018. URL <http://dblp.uni-trier.de/db/journals/corr/corr1802.html#abs-1802-06175>.
- Pang Wei Koh and Percy Liang. Understanding black-box predictions via influence functions. In Doina Precup and Yee Whye Teh (eds.), *ICML*, volume 70 of *JMLR Workshop and Conference Proceedings*, pp. 1885–1894. JMLR.org, 2017. URL <http://dblp.uni-trier.de/db/conf/icml/icml2017.html#KohL17>.

- Alex Krizhevsky. Learning multiple layers of features from tiny images. *Technical report*, 2009. URL <https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>.
- M Pawan Kumar, Benjamin Packer, and Daphne Koller. Self-Paced Learning for Latent Variable Models. In *Proc. of NIPS*, pp. 1–9, 2010.
- Y. LeCun, C. Cortes C., and C. Burges. The mnist database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>, 1999.
- Yong Jae Lee and Kristen Grauman. Learning the easy things first: Self-paced visual category discovery. In *Computer Vision and Pattern Recognition (CVPR), 2011 IEEE Conference on*, pp. 1721–1728. IEEE, 2011.
- Chunyuan Li, Heerad Farkhor, Rosanne Liu, and Jason Yosinski. Measuring the intrinsic dimension of objective landscapes. *CoRR*, abs/1804.08838, 2018. URL <http://dblp.uni-trier.de/db/journals/corr/corr1804.html#abs-1804-08838>.
- Michael McCloskey and Neal J Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. In *Psychology of learning and motivation*, volume 24, pp. 109–165. Elsevier, 1989.
- Behnam Neyshabur, Ryota Tomioka, and Nathan Srebro. In search of the real inductive bias: On the role of implicit regularization in deep learning. *CoRR*, abs/1412.6614, 2014. URL <http://dblp.uni-trier.de/db/journals/corr/corr1412.html#NeyshaburTS14>.
- Guillermo Valle Perez, Chico Q. Camargo, and Ard A. Louis. Deep learning generalizes because the parameter-function map is biased towards simple functions. *CoRR*, abs/1805.08522, 2018. URL <http://dblp.uni-trier.de/db/journals/corr/corr1805.html#abs-1805-08522>.
- Sachin Ravi and Hugo Larochelle. Optimization as a model for few-shot learning. In *Proc. of ICLR*, 2017.
- Hippolyt Ritter, Aleksandar Botev, and David Barber. Online Structured Laplace Approximations For Overcoming Catastrophic Forgetting. *arxiv preprint arxiv: 1805.07810*, 2018. URL <http://arxiv.org/abs/1805.07810>.
- Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized experience replay. *arXiv preprint arXiv:1511.05952*, 2015.
- Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, and Nathan Srebro. The Implicit Bias of Gradient Descent on Separable Data. *arXiv preprint arxiv:1710.10345*, 2017. URL <http://arxiv.org/abs/1710.10345>.
- Sainbayar Sukhbaatar, Joan Bruna, Manohar Paluri, Lubomir Bourdev, and Rob Fergus. Training convolutional networks with noisy labels. *arXiv preprint arXiv:1406.2080*, 2014.
- R. Tachet, M. Pezeshki, S. Shabanian, A. Courville, and Y. Bengio. On the learning dynamics of deep neural networks. *arxiv preprint arxiv:1809.06848*, 2018. doi: arXiv:1809.06848v1. URL <https://arxiv.org/abs/1809.06848>.
- Huan Wang, Nitish Shirish Keskar, Caiming Xiong, and Richard Socher. Identifying Generalization Properties in Neural Networks. *arXiv preprint arXiv:1809.07402*, pp. 1–23, 2018. doi: arXiv:1809.07402v1. URL <http://arxiv.org/abs/1809.07402>.
- Tengyu Xu, Yi Zhou, Kaiyi Ji, and Yingbin Liang. Convergence of sgd in learning relu models with separable data. *CoRR*, abs/1806.04339, 2018. URL <http://dblp.uni-trier.de/db/journals/corr/corr1806.html#abs-1806-04339>.
- Sergey Zagoruyko and Nikos Komodakis. Wide residual networks, 2016. URL <http://arxiv.org/abs/1605.07146>. cite arxiv:1605.07146.

Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. Understanding deep learning requires rethinking generalization. *arXiv preprint arXiv:1611.03530*, 2016.

Peilin Zhao and Tong Zhang. Stochastic Optimization with Importance Sampling for Regularized Loss Minimization. In *Proc. of ICML*, 2015.