




Figure 18: Generalization performance on *CIFAR-100* of ResNet18 where increasingly larger subsets of the training set are removed (mean $\pm$ std error of 5 seeds). When the removed examples are selected at random, performance drops faster. Selecting the examples according to our ordering reduces the training set without affecting generalization.